Developed a machine learning model to predict future stock prices for a specific company or a set of companies.

Data:

Historical stock price data (including open, close, high, low, volume)

Methodology:

Data pre-processing: Clean, transform

Model selection: Experiment with different machine learning models, such as: Linear Regression

Model training and evaluation: Train and evaluate the models using techniques like cross-validation and metrics like mean squared error (MSE), mean absolute error (MAE)

Back testing: Evaluate the model's performance on historical data not used for training. Deliverables:

A trained and evaluated machine learning model for stock price prediction.

Visualization of the predicted vs. actual stock prices.

Considerations:

Stock price prediction is inherently challenging due to market volatility and uncertainty. No model can perfectly predict future prices. The project aims to improve prediction accuracy compared to simple baselines. Ethical considerations and responsible use of the model are important.